

# Cuallee: A Python package for data quality across multiple DataFrame APIs

Herminio Vazquez <sup>1\*</sup> and Virginie Grosboillot <sup>2\*</sup>

<sup>1</sup> Independent Researcher, Mexico <sup>2</sup> Swiss Federal Institute of Technology (ETH) \* These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

## Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

## License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))

## Summary

In today's world, where vast amounts of data are generated and collected daily, and where data heavily influence business, political, and societal decisions, it's crucial to evaluate the quality of the data used for analysis, decision-making, and reporting. This involves understanding how reliable and trustworthy the data are. To address this need, we've created cuallee, a Python package for assessing data quality. cuallee is designed to be dataframe-agnostic, offering an intuitive and user-friendly API for describing checks across the most popular dataframe implementations such as PySpark, Pandas, Snowpark, Polars, DuckDB and BigQuery. Currently, cuallee offers over 50 checks to help users evaluate the quality of their data.

## Statement of need

For data engineers and data scientists, maintaining a consistent workflow involves operating in hybrid environments, where they develop locally before transitioning data pipelines and analysis to cloud-based environments. Whilst working in local environments typically allows them to fit data sets in memory, moving workloads to cloud environments involve operating with full scale data that requires a different computing framework, i.e. distributed computing, parallelization, and horizontal scaling.

This shift in computing frameworks requires the adoption of testing strategies that can accommodate testing activities in both local and remote environments, without the need to rewrite test scenarios or employ different testing approaches for assessing various quality dimensions of the data.

An additional argument is related to the rapid evolution of the data ecosystem. Organizations and data teams are constantly seeking ways to improve, whether through cost-effective solutions or by integrating new capabilities into their data operations. However, this pursuit presents new challenges when migrating workloads from one technology to another. As information technology and data strategies become more resilient against vendor lock-ins, they turn to technologies that enable seamless operation across platforms, avoiding the chaos of fully re-implementing data products. In essence, no data testing strategy needs to be rewritten or reformulated due to platform changes.

A last argument is the need for such a quality tool, is the desire of moving quality procedures to the earliest phases of the data product development life-cycle. Whether in industry or academia, the reduction of time allocated for quality activities is unfortunately like a norm, due to the predominant focus on functional aspects. Enabling a declarative, intuitive and flexible programming interface to data quality, allows teams to embed quality into their development, adopting a mindset of building quality in, as opposed to testing quality out.

## 40 Checks

41 (Binney & Tremaine, 2008) is the reference for the checks

Check	Description	Data Type
is_complete	Zero nulls	agnostic
is_unique	Zero duplicates	agnostic
is_primary_key	Zero duplicates	agnostic
are_complete	Zero nulls on group of columns	agnostic
are_unique	Composite primary key check	agnostic
is_composite_key	Zero duplicates on multiple columns	agnostic
is_greater_than	col > x	numeric
is_positive	col > 0	numeric
is_negative	col < 0	numeric
is_greater_or_equal_than	col >= x	numeric
is_less_than	col < x	numeric
is_less_or_equal_than	col <= x	numeric
is_equal_than	col == x	numeric
is_contained_in	col in [a, b, c, ...]	agnostic
is_in	Alias of is_contained_in	agnostic
not_contained_in	col not in [a, b, c, ...]	agnostic
not_in	Alias of not_contained_in	agnostic
is_between	a <= col <= b	numeric, date
has_pattern	Matching a pattern defined as a regex	string
is_legit	String not null & not empty ^\S\$	string
has_min	min(col) == x	numeric
has_max	max(col) == x	numeric
has_std	$\sigma(\text{col}) == x$	numeric
has_mean	$\mu(\text{col}) == x$	numeric
has_sum	$\Sigma(\text{col}) == x$	numeric
has_percentile	$\%(\text{col}) == x$	numeric
has_cardinality	count(distinct(col)) == x	agnostic
has_max_by	A utility predicate for max(col_a) == x for max(col_b)	agnostic
has_min_by	A utility predicate for min(col_a) == x for min(col_b)	agnostic
has_correlation	Finds correlation between 0..1 on corr(col_a, col_b)	numeric
has_entropy	Calculates the entropy of a column entropy(col) == x for classification problems	numeric
is_inside_interquartile_range	Verifies column values reside inside limits of interquartile range Q1 <= col <= Q3 used on anomalies.	numeric
is_in_millions	col >= 1e6	numeric
is_in_billions	col >= 1e9	numeric
is_t_minus_1	For date fields confirms 1 day ago t-1	date
is_t_minus_2	For date fields confirms 2 days ago t-2	date
is_t_minus_3	For date fields confirms 3 days ago t-3	date
is_t_minus_n	For date fields confirms n days ago t-n	date
is_today	For date fields confirms day is current date t-0	date
is_yesterday	For date fields confirms 1 day ago t-1	date

Check	Description	DataType
is_on_weekday	For date fields confirms day is between Mon-Fri	<i>date</i>
is_on_weekend	For date fields confirms day is between Sat-Sun	<i>date</i>
is_on_monday	For date fields confirms day is Mon	<i>date</i>
is_on_tuesday	For date fields confirms day is Tue	<i>date</i>
is_on_wednesday	For date fields confirms day is Wed	<i>date</i>
is_on_thursday	For date fields confirms day is Thu	<i>date</i>
is_on_friday	For date fields confirms day is Fri	<i>date</i>
is_on_saturday	For date fields confirms day is Sat	<i>date</i>
is_on_sunday	For date fields confirms day is Sun	<i>date</i>
is_on_schedule	For date fields confirms time windows i.e. 9:00 - 17:00	<i>timestamp</i>
is_daily	Can verify daily continuity on date fields by default. [2,3,4,5,6] which represents Mon-Fri in PySpark. However new schedules can be used for custom date continuity	<i>date</i>
has_workflow	Adjacency matrix validation on 3-column graph, based on group, event, order columns.	<i>agnostic</i>
satisfies	An open SQL expression builder to construct custom checks	<i>agnostic</i>
validate	The ultimate transformation of a check with a dataframe input for validation	<i>agnostic</i>

## Controls

This are the controls

## References

- Binney, J., & Tremaine, S. (2008). *Galactic Dynamics: Second Edition*. Princeton University Press. <http://adsabs.harvard.edu/abs/2008gady.book.....B>