

cuallee: A python package for data quality checks across multiple DataFrame APIs

Herminio Vazquez ^{1*} and Virginie Grosboillot ^{2*}

¹ Independent Researcher, Mexico ² Swiss Federal Institute of Technology (ETH) * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

In today's world, where vast amounts of data are generated and collected daily, and where data heavily influence business, political, and societal decisions, it is crucial to evaluate the quality of the data used for analysis, decision-making, and reporting. This involves understanding how reliable and trustworthy the data are. To address this need, we have created cuallee, a Python package for assessing data quality. cuallee is designed to be dataframe-agnostic, offering an intuitive and user-friendly API for describing checks across the most popular dataframe implementations such as PySpark, Pandas, Snowpark, Polars, DuckDB and BigQuery. Currently, cuallee offers over 50 checks to help users evaluate the quality of their data.

Statement of need

For data engineers and data scientists, maintaining a consistent workflow involves operating in hybrid environments, where they develop locally before transitioning data pipelines and analyses to cloud-based environments. Whilst working in local environments typically allows them to fit data sets in memory, moving workloads to cloud environments involve operating with full scale data that requires a different computing framework ([Schelter et al., 2018](#)), i.e. distributed computing, parallelization, and horizontal scaling. cuallee accomodates the testing activities required by this shift in computing frameworks, in both local and remote environments, without the need to rewrite test scenarios or employ different testing approaches for assessing various quality dimensions of the data ([Fadlallah et al., 2023b](#)).

An additional argument is related to the rapid evolution of the data ecosystem ([Fadlallah et al., 2023a](#)). Organizations and data teams are constantly seeking ways to improve, whether through cost-effective solutions or by integrating new capabilities into their data operations. However, this pursuit presents new challenges when migrating workloads from one technology to another. As information technology and data strategies become more resilient against vendor lock-ins, they turn to technologies that enable seamless operation across platforms, avoiding the chaos of fully re-implementing data products. In essence, with cuallee no data testing strategy needs to be rewritten or reformulated due to platform changes.

One last argument in favor of using a quality tool such as cuallee is the need to integrate quality procedures into the early stages of data product development. Whether in industry or academia, there is often a tendency to prioritize functional aspects over quality, leading to less time being dedicated to quality activities. By providing a clear, easy-to-use, and adaptable programming interface for data quality, teams can incorporate quality into their development process, promoting a proactive approach of building quality in rather than relying solely on testing to ensure quality.

Data Quality Frameworks

Data platforms have diversified from file systems and relational databases, to full ecosystems including the concept of data lakes (Dumke et al., 2020). Modern platforms host a variety of data formats, beyond traditional tabular data, including semi-structured like JSON (Pezoa et al., 2016) or unstructured like audio or images.

Operating with modern data platforms, requires a versatile data processing framework capable to handle structured and unstructured data, support data operations in various programming languages, fulfill the imperative and declarative form to data operations from practitioners and do it reliably for any size of data. Apache Spark (Armbrust et al., 2015) represents an exemplar framework due to the wide range of data processing capabilities—batch processing, real-time streaming, machine learning, and graph processing—within a unified framework commended and adopted (O’Reilly Media, 2023) by the data industry.

cuallee is powered by native data engines, including Apache Spark, and offers a robust structure that can be extended to new engines with fully open-source implementation guidelines and rigorous testing. pydeequ (Schelter et al., 2018) is a pioneer in large-scale data quality frameworks and is fully open-source. However, its adoption is limited due to the smaller community of developers proficient in the scala programming language.

On the other hand, great-expectations and soda are commercial options that require registration and issuing of keys for cloud reporting capabilities.

cuallee provides a fully open-source data quality framework designed for both academia and industry practitioners, offering unparalleled performance compared to the aforementioned alternatives.

Performance Benchmark

A reproducible performance benchmark is available in the code repository (Vazquez, 2024). It consists in 38 checks over an open sourced data set (New York City Taxi and Limousine Commission, 2024) made of 19.8 million rows. The validation performs 19 checks for completeness and 19 checks for uniqueness for each column of the dataset.

The following table (Table 1) provides a summary of the performance comparisson:

Table 1: Performance comparisson on popular data quality frameworks

| Framework | Definitions | Time |
|-----------------------------|-------------|------|
| great_expectations==0.18.13 | python | 66s |
| soda==1.4.10 | yaml | 43s |
| pydeequ==1.3.0 | python | 11s |
| cuallee==0.10.3 | python | 7s |

Methods

cuallee employs a heuristic-based approach to define quality rules for each dataset. This prevents the inadvertent duplication of quality predicates, thus reducing the likelihood of human error in defining rules with identical predicates. Several studies have been conducted on the efficiency of these rules, including auto-validation and auto-definition using profilers (Tu et al., 2023).

74 Checks

75 In *cuallee*, checks serve as the fundamental concept. These checks (Table 2) are implemented
 76 by *rules*, which specify *quality predicates*. These predicates, when aggregated, form the criteria
 77 used to evaluate the quality of a dataset. Efforts to establish a universal quality metric
 78 (Pleimling et al., 2022) typically involve using statistics and combining dimensions to derive a
 79 single reference value that encapsulates overall quality attributes.

Table 2: List and description of the currently available checks

| Check | Description | DataType |
|--------------------------|---|---------------|
| is_complete | Zero nulls | agnostic |
| is_unique | Zero duplicates | agnostic |
| is_primary_key | Zero duplicates | agnostic |
| are_complete | Zero nulls on group of columns | agnostic |
| are_unique | Composite primary key check | agnostic |
| is_composite_key | Zero duplicates on multiple columns | agnostic |
| is_greater_than | col > x | numeric |
| is_positive | col > 0 | numeric |
| is_negative | col < 0 | numeric |
| is_greater_or_equal_than | col >= x | numeric |
| is_less_than | col < x | numeric |
| is_less_or_equal_than | col <= x | numeric |
| is_equal_than | col == x | numeric |
| is_contained_in | col in [a, b, c, ...] | agnostic |
| is_in | Alias of is_contained_in | agnostic |
| not_contained_in | col not in [a, b, c, ...] | agnostic |
| not_in | Alias of not_contained_in | agnostic |
| is_between | a <= col <= b | numeric, date |
| has_pattern | Matching a pattern defined as a regex | string |
| is_legit | String not null & not empty $^{\wedge}\backslash S\$$ | string |
| has_min | min(col) == x | numeric |
| has_max | max(col) == x | numeric |
| has_std | $\sigma(\text{col}) == x$ | numeric |
| has_mean | $\mu(\text{col}) == x$ | numeric |
| has_sum | $\Sigma(\text{col}) == x$ | numeric |
| has_percentile | $\%(\text{col}) == x$ | numeric |
| has_cardinality | count(distinct(col)) == x | agnostic |
| has_max_by | A utility predicate for max(col_a) == x for max(col_b) | agnostic |
| has_min_by | A utility predicate for min(col_a) == x for min(col_b) | agnostic |
| has_correlation | Finds correlation between 0..1 on corr(col_a, col_b) | numeric |
| has_entropy | Calculates the entropy of a column entropy(col) == x for classification problems | numeric |
| is_inside_iqr | Verifies column values reside inside limits of interquartile range Q1 <= col <= Q3 used on anomalies. | numeric |
| is_in_millions | col >= 1e6 | numeric |
| is_in_billions | col >= 1e9 | numeric |
| is_t_minus_1 | For date fields confirms 1 day ago t-1 | date |
| is_t_minus_2 | For date fields confirms 2 days ago t-2 | date |

| Check | Description | DataType |
|--------------------------|--|-----------|
| is_t_minus_3 | For date fields confirms 3 days ago t-3 | date |
| is_t_minus_n | For date fields confirms n days ago t-n | date |
| is_today | For date fields confirms day is current date t-0 | date |
| is_yesterday | For date fields confirms 1 day ago t-1 | date |
| is_on_weekday | For date fields confirms day is between Mon-Fri | date |
| is_on_weekend | For date fields confirms day is between Sat-Sun | date |
| is_on_monday | For date fields confirms day is Mon | date |
| is_on_tuesday | For date fields confirms day is Tue | date |
| is_on_wednesday | For date fields confirms day is Wed | date |
| is_on_thursday | For date fields confirms day is Thu | date |
| is_on_friday | For date fields confirms day is Fri | date |
| is_on_saturday | For date fields confirms day is Sat | date |
| is_on_sunday | For date fields confirms day is Sun | date |
| is_on_schedule | For date fields confirms time windows i.e. 9:00 - 17:00 | timestamp |
| is_daily | Can verify daily continuity on date fields by default. [2,3,4,5,6] which represents Mon-Fri in PySpark. However new schedules can be used for custom date continuity | date |
| has_workflow | Adjacency matrix validation on 3-column graph, based on group, event, order columns. | agnostic |
| satisfies | An open SQL expression builder to construct custom checks | agnostic |
| validate | The ultimate transformation of a check with a dataframe input for validation | agnostic |
| iso.iso_4217 | currency compliant ccy | string |
| iso.iso_3166 | country compliant country | string |
| Control.completeness | Zero nulls all columns | agnostic |
| Control.percentage_fill | % rows not empty | agnostic |
| Control.percentage_empty | % rows empty | agnostic |

References

- Armbrust, M., Xin, R. S., Lian, C., Huai, Y., Liu, D., Bradley, J. K., Meng, X., Kaftan, T., Franklin, M. J., Ghodsi, A., & Zaharia, M. (2015). Spark SQL: Relational data processing in spark. *Proceedings of the 2015 ACM SIGMOD International Conference on Management of Data*, 1383–1394. <https://doi.org/10.1145/2723372.2742797>
- Dumke, A. R., Parchmann, A., Schmid, S., & Hauswirth, M. (2020). Toward data lakes as central building blocks for data management and analysis. *Frontiers in Big Data*, 3, 564115. <https://doi.org/10.3389/fdata.2020.564115>
- Fadlallah, H., Kilany, R., Dhayne, H., El Haddad, R., Haque, R., Taher, Y., & Jaber, A. (2023a). BIGQA: Declarative big data quality assessment. *Journal of Data and Information Quality*, 15. <https://doi.org/10.1145/3603706>
- Fadlallah, H., Kilany, R., Dhayne, H., El Haddad, R., Haque, R., Taher, Y., & Jaber, A. (2023b). Context-aware big data quality assessment: A scoping review. *Journal of Data*

- 93 and Information Quality, 15. <https://doi.org/10.1145/3603707>
- 94 New York City Taxi and Limousine Commission. (2024). *TLC trip record data*. <https://www.nyc.gov/site/tlc/about/tlc-trip-record-data.page>
- 95
- 96 O'Reilly Media. (2023). *Technology trends for 2023*. [https://www.oreilly.com/radar/](https://www.oreilly.com/radar/technology-trends-for-2023/)
- 97 [technology-trends-for-2023/](https://www.oreilly.com/radar/technology-trends-for-2023/)
- 98 Pezoa, F., Reutter, J., Suarez, F., Ugarte, M., & Vrgoč, D. (2016). *Foundations of JSON*
- 99 *schema*. 263–273. <https://doi.org/10.1145/2872427.2883029>
- 100 Pleimling, X., Shah, V., & Lourentzou, I. (2022). [Data] quality lies in the eyes of the beholder.
- 101 *Proceedings of the 15th International Conference on PErvasive Technologies Related to*
- 102 *Assistive Environments*, 118–124. <https://doi.org/10.1145/3529190.3529222>
- 103 Schelter, S., Lange, D., Schmidt, P., Celikel, M., Biessmann, F., & Grafberger, A. (2018).
- 104 Automating large-scale data quality verification. *Proc. VLDB Endow.*, 11(12), 1781–1794.
- 105 <https://doi.org/10.14778/3229863.3229867>
- 106 Tu, D., He, Y., Cui, W., Ge, S., Zhang, H., Han, S., Zhang, D., & Chaudhuri, S. (2023).
- 107 Auto-validate by-history: Auto-program data quality constraints to validate recurring data
- 108 pipelines. *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and*
- 109 *Data Mining*, 4991–5003. <https://doi.org/10.1145/3580305.3599776>
- 110 Vazquez, H. (2024). *Cuallee: Performance tests*. [https://github.com/canimus/cuallee/tree/](https://github.com/canimus/cuallee/tree/main/test/performance)
- 111 [main/test/performance](https://github.com/canimus/cuallee/tree/main/test/performance)

DRAFT